

Testing Plan: AI-Based Anomaly Detection in Spacecraft Telemetry Data

Objective

To Validate the AI-based anomaly detection system using Isolation Forest and LSTM autoencoder for spacecraft telemetry data, ensuring >90% precision, >95% recall, <1 second latency, and robustness across edge cases, as outlined in the solution plan by Jazmine Brown and Troy Nsofar.

Testing Goals

- Confirm the system detects anomalies with high precision (>90%) and recall (>95%).
- Ensure robustness against noisy or incomplete telemetry data.
- Verify real-time performance (<1 second latency per data window).
- Address interpretability by providing clear, actionable alerts.

Test Data

- **Sources:**
 - **NASA SMAP Dataset:**- 1 million data points 10 sensors, 10 Hz sampling 3 hours of telemetry 100 MB.
 - **Curiosity Dataset:** 500,000 data points 5 sensors, 5 Hz sampling, 5 hours 50 MB.
 - **Synthetic Data:** 100,000 data points with 20% injected anomalies.
- **Characteristics:**
 - Normal data: 80% of dataset (e.g., 10,000 60-second windows at 10Hz).
 - Anomalous data: 20% with labeled anomalies (e.g., 2,000 windows with spikes or drops).
 - Edge cases: 10% missing values, Gaussian noise (SNR=10), or extreme readings ($\pm 5\sigma$).
- **Data Split:**
 - Training: 70% (normal data only for LSTM autoencoder; mixed for Isolation Forest).
 - Validation: 15% (mixed normal and anomalous for threshold tuning).
 - Testing: 15% (mixed normal and anomalous for final evaluation).

Testing Methods

1. Unit Testing:

- **Preprocessing:** Verify data cleaning (no NaN values), normalization (values in $[0, 1]$), and windowing (correct 60-second segments).
- **Isolation Forest:** Ensure model outputs anomaly scores (0 to 1) for static data.
- **LSTM Autoencoder:** Confirm reconstruction errors are generated for time-series inputs.
- **Alert System:** Test API delivers JSON alerts with correct format (e.g., timestamp, sensor, score).

2. Integration Testing:

- Test end-to-end pipeline: Telemetry input → preprocessing → model (Isolation Forest or LSTM) → alert output.
- Verify API integration with a mock mission control dashboard.
- Evaluate edge cases missing data, noise, corrupted inputs with >85% recall target.

3. Performance Testing:

- Measure latency: Process 1,000 data windows; target <1 second per window.
- Test scalability: Run on 1 million data points to ensure no crashes.

4. Robustness Testing:

- Test edge cases: Missing data (10% of points), noisy data (SNR=10), or corrupted inputs.
- Verify model performance (recall >85%) under adverse conditions.

Evaluation Metrics

• Primary Metrics:

- **Precision:** Percentage of flagged anomalies that are true (target: >90%).
- **Recall:** Percentage of true anomalies detected (target: >95%).
- **F1-Score:** Harmonic mean of precision and recall (target: >92%).

• Secondary Metrics:

- **False Positive Rate:** Normal data flagged as anomalous target: <5%.
- **Latency:** Processing time per window target: <1 second per 1,000 windows.
- **Threshold Tuning:**
 - Adjust Isolation Forest contamination (0.1 to 0.15) and LSTM reconstruction error threshold (0.05 to 0.1) to optimize metrics.
 - LSTM: Tune reconstruction error threshold to optimize F1-score.

Test Scenarios

Normal Operation:

- Input: Normal telemetry (e.g., stable temperature, voltage).
- Expected: No alerts; Isolation Forest scores <0.7; LSTM errors below threshold.

Known Anomaly:

- Input: Data with injected anomaly (e.g., 50°C temperature spike).
- Expected: Alert triggered with correct details (e.g., timestamp, sensor, score >0.7).

Edge Case:

- Input: Data with 10% missing values or added noise (SNR=10).
- Expected: Recall >85%; alerts remain accurate.

High Load:

- Input: Continuous 24-hour telemetry stream (e.g., 864,000 data points at 10Hz).
- Expected: System processes data in real time without crashes; latency <1 second.

Testing Process

1. Setup:

- Use simulated NASA data SMAP, Curiosity, synthetic on AWS EC2, mimicking deployment environment.

2. Execution:

- Run unit tests on preprocessing, models, and alerts.
- Conduct integration tests on the full pipeline.
- Perform performance and robustness tests with edge cases.

3. **Analysis:**

- Calculate metrics precision, recall, F1-score, false positive rate for each scenario.
- Log results in docs/test_results.md with tables and visualizations (e.g., ROC curves).

4. **Iteration:**

- If metrics miss targets precision <90%, recall <95%
 1. Retraining: Adjust hyperparameters like Isolation Forest contamination to 0.15, LSTM learning rate to .0001.
 2. Fallback Model: Use Isolation Forest Alone if LSTM latency exceeds 1 second, leveraging its faster inference.
 3. Data Adjustment: Increase training data by 20% or augment with synthetic normal data to improve model robustness.
 4. Retest until targets are met.

Addressing Challenges

• **Data Quality:**

- Test preprocessing robustness with noisy or incomplete data.
- Use Isolation Forest as a fallback for low-quality data scenarios.

• **Real-Time Constraints:**

- Optimize LSTM inference (e.g., model quantization) and leverage Isolation Forest's speed.
- Measure latency in high-load tests to ensure <1 second processing.

• **Model Interpretability:**

- Include visualizations (e.g., time-series plots with flagged anomalies) in alerts.
- Provide confidence scores and sensor details to aid mission control.

Documentation

- Record results in docs/test_results.md:
 - Tables of metrics (precision, recall, F1-score) per scenario.
 - Visualizations (e.g., anomaly detection plots).
 - Notes on failures and mitigations.
- Update README.md: “See docs/testing_plan.md for the testing strategy and docs/test_results.md for results.”

Expected Outcomes

- System achieves >90% precision, >95% recall, and <5% false positives.
- Processes telemetry in real time (<1 second per 1,000 windows latency).
- Robust to edge case recall, maintaining >85% recall in noisy conditions.
- Provides clear, interpretable alerts for mission control.

References

- Hundman et al., “Detecting Spacecraft Anomalies Using LSTMs and Nonparametric Dynamic Thresholding,” KDD '18: Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining, 2018.